CREATORPROOF / PHASE 2 OPTION

AI Protection for Artist Work

A concise guide to existing tools that can make AI training, style copying, and direct image editing harder - and how CreatorProof can use them responsibly.

PROTECTION +
PROVENANCE +
EVIDENCE

The honest answer

CreatorProof can generate protected copies of an artist's work. These copies can reduce the success of selected AI workflows. No current tool can guarantee that a determined attacker will never copy, train on, or recreate a public image.

What CreatorProof adds	What it does not promise
- A protected derivative for public sharing - A clean original kept private - C2PA and rights records on the final file - Evidence when misuse is detected	- Perfect protection against every AI model - Removal of an image already learned by a model - A legal decision made automatically - A replacement for artist consent or licensing

Prepared for the CreatorProof product and pitch workflow

The protection options

Each option targets a different kind of misuse. They should be offered as separate modes, because one tool cannot reliably solve every AI threat.

Glaze

Protects against: AI learning and copying an artist's visual style

How it works: Adds tiny image-specific changes that look normal to people but can confuse style-learning models.

Best use: Public portfolios and social posts where the main concern is style imitation.

Important limit: Its creators say it is not consistent protection against image-to-image editing or inpainting. [1]

PhotoGuard

Protects against: Direct AI editing, inpainting, and image-to-image changes

How it works: Adds calculated changes intended to disrupt the AI editing process.

Best use: A self-hosted CreatorProof pilot for selected Stable Diffusion editing workflows.

Important limit: Research code is not proof against every commercial or future model. Test it against supported models. [2]

Nightshade

Protects against: Scrapers that ignore opt-out instructions and add images to training data

How it works: Creates a poisoned training signal aimed at selected concepts.

Best use: An optional, high-risk mode for artists who want a stronger deterrent against non-consensual scraping.

Important limit: It is not a lock on a single image and does not stop every AI system. [3]

Mist

Protects against: Personalised AI training such as LoRA and DreamBooth

How it works: Adds adversarial changes designed to make fine-tuning produce poor results.

Best use: Research evaluation or a future licensed partnership.

Important limit: The authors say it may only raise the cost for a skilled attacker. Its repository also forbids unauthorised commercial use. [4]

The evidence check

Independent research found that common adversarial image protections can be weakened by preprocessing such as upscaling. CreatorProof should therefore describe these tools as tested risk reduction, never as an AI-proof guarantee. [5]

How CreatorProof uses them

The platform protects a copy for public use while preserving the clean original for rights evidence and licensing.

STEP	WHAT HAPPENS	WHY IT MATTERS
1	**Artist uploads**	The clean original is stored privately. No public file link is created.
2	**Artist chooses a mode**	Style protection, editing protection, or optional high-risk training deterrent.
3	**Protection worker runs**	CreatorProof generates a separate protected derivative and records the tool version.
4	**Artist reviews**	The artist compares the original and protected copy, then approves the public version.
5	**Credentials are added**	C2PA and metadata are added to the final protected file. The final file is hashed and linked to the original.
6	**Monitoring continues**	CreatorProof searches for copies, similar outputs, and possible misuse, then builds an evidence case.

Private original Rights evidence, licensing, investigation	**Protected derivative** Public portfolio, preview, social sharing	**Evidence record** Hashes, credentials, policy, detection history

Why C2PA still matters

C2PA can show where the file came from, what happened to it, and what the artist selected as their permission policy. It is a tamper-evident provenance record, not a technical barrier that stops AI training or copying. [6]

What to build first

The strongest product position is a layered system: private storage, protected derivatives, provenance, and active evidence collection.

PRIORITY	BUILD	VALUE
1	**Protected export** Generate a protected copy while preserving the clean original privately.	Most useful artist-facing feature. Creates a clear public-safe workflow.
2	**PhotoGuard pilot** Test direct editing and inpainting protection on selected open models.	Gives CreatorProof a concrete AI-editing protection demo.
3	**Glaze permission pathway** Explore approved use or artist-side local processing. Do not assume a public API exists.	Adds a recognised style-protection option without making an unsupported promise.
4	**Monitoring and evidence** Match copies and generated outputs to the artist record, then open a case.	Turns protection into a complete rights product instead of a single image filter.

Pitch-safe wording

CreatorProof creates protected, traceable versions of artwork that can make selected AI training and editing workflows less effective. It gives artists stronger control, clearer evidence, and a safer way to publish.

Sources

[1] Glaze FAQ, University of Chicago: https://glaze.cs.uchicago.edu/faq.html

[2] PhotoGuard research repository, MadryLab: https://github.com/MadryLab/photoguard

[3] Nightshade paper, IEEE Security and Privacy 2024: https://arxiv.org/abs/2310.13828

[4] Mist v2 repository and usage terms: https://github.com/psyker-team/mist-v2

[5] Hönig et al., Adversarial Perturbations Cannot Reliably Protect Artists From Generative AI: https://arxiv.org/abs/2406.12027

[6] C2PA clarification on TDM and DRM: https://c2pa.org/c2pa-clarification-to-c2pa-tdm-assertions-reference/

Decision rule: promise measurable resistance against named workflows, preserve the clean original privately, and keep monitoring for misuse. That is a credible protection product.